Temporal Fusion Transformers for Uncertainty-Quantified Stock Forecasting

Project-II (MA47202) report submitted to

Indian Institute of Technology Kharagpur

in partial fulfillment for the award of the degree of

Integrated Masters in Science

in

Mathematics and Computing

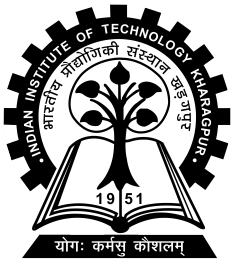
by

Sameer Pagarani

(21MA20002)

Under the supervision of

Professor Pawan Kumar



Department of Mathematics

Indian Institute of Technology Kharagpur

Spring Semester, 2025-26

May 04, 2026

DECLARATION

I certify that

(a) The work contained in this report has been done by me under the guidance of my supervisor.

(b) The work has not been submitted to any other Institute for any degree or diploma.

(c) I have conformed to the norms and guidelines given in the Ethical Code of Conduct of the Institute.

(d) Whenever I have used materials (data, theoretical analysis, figures, and text) from other sources, I have given due credit to them by citing them in the text of the thesis and giving their details in the references.

Place: Kharagpur

Date: April 25, 2026

Sameer Pagarani

(21MA20002)

DEPARTMENT OF MATHEMATICS
INDIAN INSTITUTE OF TECHNOLOGY KHARAGPUR
KHARAGPUR - 721302, INDIA



CERTIFICATE

This is to certify that the project report entitled **"Temporal Fusion Transformers for Uncertainty-Quantified Stock Forecasting"** submitted by **Sameer Pagarani** (Roll No. **21MA20002**) to Indian Institute of Technology Kharagpur towards partial fulfillment of requirements for the award of degree of **Integrated Masters in Science** in **Mathematics and Computing** is a record of bona fide work carried out by him under my supervision and guidance during Spring Semester, 2025-26.

Place: Kharagpur
Date: April 25, 2026

Professor Pawan Kumar
Department of Mathematics
Indian Institute of Technology Kharagpur
Kharagpur - 721302, India

Abstract

Equity markets are noisy, non-stationary, and prone to sudden regime shifts, making reliable short-horizon forecasting difficult yet practically valuable. This thesis presents **Temporal Fusion Transformers for Uncertainty-Quantified Stock Forecasting**, an end-to-end framework that predicts near-term price evolution while explicitly communicating uncertainty through probabilistic outputs.

Building on insights from **Project 1**—where a volatility-resilient ensemble of recurrent networks and a simple stacking strategy improved robustness—this work advances the modeling approach by adopting the **Temporal Fusion Transformer (TFT)**, a modern architecture designed to learn time-varying feature relevance and complex temporal structure. The model is trained on a curated universe of **39 U.S. equities across the semiconductor, cloud computing, and hardware sectors** using daily OHLCV data. A consistent feature pipeline combines price-based transformations, technical indicators, and market-context signals to capture both asset-specific behaviour and broader market effects.

Rather than producing a single point estimate, the proposed system generates **quantile forecasts (q10/q50/q90)**, where q50 provides the median trajectory and q10–q90 forms an uncertainty band suitable for scenario analysis and risk-aware interpretation. Forecasts are produced for short, decision-relevant horizons of **1, 3, and 5 trading days**. Held-out evaluation across the trained universe indicates low relative error (MAPE typically in the $\sim$2–4% range across horizons), supporting the effectiveness of TFT-style probabilistic forecasting for short-term equity prediction.

For reproducibility and controlled demonstration, the trained model is integrated into a password-protected web interface that serves forecasts without exposing project artifacts.

Keywords: Temporal Fusion Transformer, Stock Forecasting, Financial Time Series, Probabilistic Forecasting, Quantile Regression, Uncertainty Quantification, Multi-Horizon Prediction, Feature Engineering.

Contents

Declaration i

Certificate ii

Abstract iii

1 Introduction 1

1.1 Research Context & Motivation . . . 1

1.2 From Project 1 to Project 2: Why Move Beyond Ensembles . . . 1

1.3 Uncertainty-Quantified Forecasting with Temporal Fusion Transformers . . . 2

1.4 Key Contributions . . . 2

2 Literature Survey 3

2.1 The Statistical and Recurrent Era . . . 3

2.2 The Transformer Revolution (2021–2023) . . . 3

2.3 Frontier: Time-Series Foundation Models (2024–2026) . . . 4

2.4 Why TFT? (Conclusion) . . . 4

3 Data Collection, Preprocessing & Feature Engineering 5

3.1 Data Acquisition Strategy . . . 5

3.1.1 Primary and Fallback Sources . . . 5

3.1.2 The "Trained Universe" . . . 5

3.2 Preprocessing and Schema Standardization . . . 6

3.3 Feature Engineering . . . 6

3.3.1 1. Target Transformation (Log Returns) . . . 6

3.3.2 2. Technical Indicators (Momentum & Volatility) . . . 6

3.3.3 3. Market Context (The "Global" Signal) . . . 7

3.4 Sequence Construction for TFT . . . 7

3.4.1 Time Indexing and Grouping . . . 7

3.4.2 Encoder-Decoder Split . . . 7

4 Proposed System & Methodology **8**
4.1 Why TFT for Real Markets? . . . 8
4.2 Forecasting Objective . . . 8
4.3 Data Representation and Input Signals . . . 8
4.4 Temporal Fusion Transformer (TFT) Backbone . . . 9
4.5 Multi-Horizon Probabilistic Forecasting . . . 10
4.6 Training Strategy and Generalization Across Stocks . . . 10
4.7 Deployment Architecture (Streamlit on Hugging Face) . . . 10

5 Temporal Fusion Transformer (TFT) Architecture **11**
5.1 Problem Formulation . . . 11
5.2 Input Decomposition . . . 11
5.3 Core Architectural Components . . . 12
5.4 Inference Algorithm . . . 13
5.5 Quantile Loss Function . . . 13

6 Deployment & Web Application **14**
6.1 System Architecture . . . 14
6.2 Execution Pipeline: From Click to Forecast . . . 14
6.3 User Interface . . . 15
6.3.1 Interactive Forecasting . . . 15
6.3.2 Documentation & Reporting . . . 16
6.4 Engineering for Reproducibility . . . 17

7 Results & Evaluation **18**
7.1 Evaluation Protocol . . . 18
7.2 Metrics Definitions . . . 18
7.3 Comprehensive Results . . . 19
7.3.1 Day 1 Horizon (Short-Term) . . . 19
7.3.2 Day 3 Horizon (Medium-Term) . . . 19
7.3.3 Day 5 Horizon (Long-Term) . . . 20
7.4 Qualitative Analysis . . . 20
7.5 Conclusion . . . 20

8 Conclusion & Future Direction **23**

A Trained Universe List (39 Tickers) **25**

References **26**

List of Figures

3.1 The **data pipeline** transforming raw, noisy inputs into an aligned, context-aware tensor formatted for the TFT. 6

4.1 **Conceptual Pipeline Overview.** The end-to-end forecasting flow: heterogeneous inputs → variable selection/gating → sequence modeling and attention → multi-horizon probabilistic outputs. 9

5.1 **TFT Architecture Overview.** The data flows from left to right: (1) Inputs are weighted by Variable Selection Networks, (2) Non-linearities are handled by Gated Residual Networks, (3) Local dynamics are captured by LSTM layers, and (4) Long-range dependencies are modeled via Interpretable Multi-Head Attention. 12

6.1 **Deployment Architecture.** The system integrates live data retrieval with the pre-trained TFT checkpoint to output q_{10}, q_{50}, q_{90} probability paths, rendered as interactive plots. 14

6.2 **Live Forecast Dashboard.** This allows users to select the target asset, horizon and other settings - such as history window, force live fetch and bias calibrate. 15

6.3 **Embedded User Guide.** A dedicated 'FAQ' tab in the app. 16

6.4 **Project Documentation Access.** The 'Report' tab provides a direct download link for the full Technical Project Report. This allows users to review the underlying research and model architecture alongside the live demonstration. 16

7.1 **Stable Market Example.** The median prediction tracks price closely; the 80% band is narrow, indicating high certainty. 21

7.2 **Volatile Market Example.** The model widens the interval during turbulence, effectively capturing risk. 22

Chapter 1

Introduction

1.1 Research Context & Motivation

Equity markets are complex, stochastic systems driven by a multitude of factors: prices react not only to fundamental valuation, but also to macroeconomic shocks, liquidity constraints, and rapid shifts in investor sentiment. For quantitative practitioners, the challenge lies not merely in predicting directional movement, but in decision-making under uncertainty—small timing errors can compound into significant drawdowns, particularly during volatile regimes. This makes short-horizon forecasting (1–5 days) simultaneously critical and notoriously difficult: the signal-to-noise ratio is low, return distributions are heavy-tailed, and market dynamics exhibit non-stationarity.

Classical statistical models often struggle in this domain due to rigid assumptions of linearity or homoscedasticity. In contrast, real financial time series are characterized by nonlinear interactions and abrupt regime changes where historical patterns may suddenly cease to hold. Consequently, a single point estimate is rarely sufficient for robust decision-making. A practical forecasting system must instead communicate confidence and risk, providing a probabilistic range of outcomes rather than a false sense of certainty.

This project is motivated by this necessity: to build a forecasting framework that is accurate when signal is strong, cautious when noise is high, and explicitly quantified in its uncertainty. The objective is to construct a rigorous, reproducible system that extracts learnable temporal structure from financial data while remaining interpretable under volatility.

1.2 From Project 1 to Project 2: Why Move Beyond Ensembles

This work serves as a natural evolution of **Project 1**, which explored a "Volatility-Resilient Ensemble" of recurrent networks (LSTM, GRU, RNN). That project utilized a linear stacking strategy to combine complementary architectural biases, establishing a strong baseline for stability. However, the ensemble approach highlighted a critical limitation: static stacking mechanisms struggle to adapt dynamically when the relevance of input features shifts over time.

Project 2 addresses this limitation by adopting the **Temporal Fusion Transformer (TFT)**. Unlike standard recurrence, the TFT employs gated feature selection and attention mechanisms to dynamically determine which signals are relevant at each specific time step. This capability is uniquely aligned with financial markets, where the driver of price action can shift rapidly between technical momentum, volatility clusters, or broad market correlation. If Project 1 demonstrated that architectural diversity reduces variance, Project 2 demonstrates that context-aware attention improves adaptability within a single unified architecture.

1.3 Uncertainty-Quantified Forecasting with Temporal Fusion Transformers

The central methodological contribution of Project 2 is **uncertainty-aware, multi-horizon forecasting**. Rather than outputting a deterministic trajectory, the model generates **quantile forecasts** (q10/q50/q90). Here, the q50 represents the median prediction, while the interval between q10 and q90 forms a calibrated uncertainty band. This is operationally valuable: in stable regimes, the band tightens, signaling high confidence; in turbulent markets, the band widens, effectively warning the user of heightened risk.

To ensure the study is grounded in realistic volatility, the model is trained and evaluated on a curated universe of **39 U.S. equities across the semiconductor, cloud computing, and hardware sectors**. The forecasting task targets short decision windows of **1, 3, and 5 trading days**. The input pipeline constructs a consistent feature frame from daily OHLCV data, integrating price-based transformations, technical indicators, and broad market-context signals to separate idiosyncratic alpha from systematic beta.

Finally, while the thesis emphasizes modeling and evaluation, the work includes a lightweight deployment layer implemented as a secure web application. This ensures reproducibility and accessibility, allowing the trained checkpoint to be tested on out-of-sample data and demonstrated to users through a controlled interface without exposing internal project artifacts.

1.4 Key Contributions

Within the scope of this thesis, the primary contributions are:

- A **TFT-based forecasting framework** for short-horizon equity prediction on a curated 39-ticker universe, utilizing variable selection networks for feature importance.
- **Probabilistic (quantile) forecasting** that provides calibrated uncertainty bands alongside median estimates, enabling risk-aware interpretation.
- A **reproducible end-to-end pipeline** linking rigorous data preparation, multi-horizon evaluation, and a secure demonstration interface.

Chapter 2

Literature Survey

Financial forecasting sits at the intersection of stochastic processes and representation learning. The field has evolved from rigid statistical assumptions to data-driven deep learning, and most recently, to large-scale foundation models. This chapter briefly traces that trajectory to contextualize why the **Temporal Fusion Transformer (TFT)** remains the optimal design choice for this thesis.

2.1 The Statistical and Recurrent Era

For decades, the standard for modeling financial time series relied on linear autoregressive frameworks (ARIMA) and heteroskedasticity-aware variance models (GARCH) [1]. While mathematically elegant, these models impose strong stationarity assumptions that rarely hold in modern, regime-switching markets.

The emergence of Deep Learning brought Recurrent Neural Networks (RNNs) and LSTMs to the forefront. Fischer and Krauss [2] demonstrated that LSTMs could capture non-linear dependencies that statistical models missed. However, one can spot the bottleneck immediately: RNNs process data sequentially. This results in the "vanishing gradient" problem over long sequences, making them poor at capturing long-range market cycles or delayed reactions to macro shocks.

2.2 The Transformer Revolution (2021–2023)

The paradigm shifted with the introduction of the Transformer. By replacing recurrence with *Self-Attention*, models could attend to relevant historical events regardless of distance. Initial adaptations like *Informer* [3] and *Autoformer* [4] solved the computational complexity of standard attention ($\mathcal{O}(L^2)$) using sparse approximations and decomposition mechanisms, establishing a new SOTA for long-sequence forecasting.

However, generic Transformers often treat inputs as black boxes. This led to the **Temporal Fusion Transformer (TFT)** [5], which introduced an inductive bias specifically for multi-horizon forecasting:

1. **Gating Mechanisms (GRN):** To suppress noise and select only relevant features (essential for

low signal-to-noise ratio financial data).

2. **Quantile Regression:** Instead of a single mean prediction, TFT optimizes for the full conditional distribution (P_{10}, P_{50}, P_{90}), making it natively probabilistic.

2.3 Frontier: Time-Series Foundation Models (2024–2026)

The current "bleeding edge" involves adapting Large Language Models (LLMs) for time series. Models like **TimeGPT** [6] and Amazon's **Chronos** [7] tokenize continuous values to treat forecasting as a next-token classification task. *PatchTST* [8] further refined this by segmenting series into "patches" to capture local semantics.

2.4 Why TFT? (Conclusion)

While Foundation Models are impressive, they are often computationally heavy and lack interpretability. In contrast, RNNs are efficient but lack global context. The **TFT** sits in the "Goldilocks" zone for financial engineering: it is lighter than an LLM, more powerful than an LSTM, and crucially, it provides *interpretability*—allowing us to verify which market signals drive the forecast.

Chapter 3

Data Collection, Preprocessing & Feature Engineering

The quality of a financial forecasting model is bounded by its input data. Unlike computer vision or NLP tasks where raw data often suffices, financial time series require rigorous preprocessing to handle non-stationarity, noise, and regime shifts.

This chapter details the end-to-end data pipeline developed for this thesis. We transition from a univariate approach (Project 1) to a multivariate, context-aware framework (Project 2), specifically designed to leverage the variable selection capabilities of the Temporal Fusion Transformer (TFT).

3.1 Data Acquisition Strategy

The system relies on daily Open-High-Low-Close-Volume (OHLCV) data. To ensure the deployed application remains operational in production, we implemented a redundant data retrieval strategy.

3.1.1 Primary and Fallback Sources

- **Primary Source (Yahoo Finance):** We utilize the `yfinance` API to fetch adjusted historical data. This serves as the default provider due to its broad coverage and split adjustments.
- **Fallback Source (Stooq):** Financial APIs are prone to rate limits. We engineered an automatic fallback that queries the Stooq database if the primary source fails, ensuring high availability.

3.1.2 The "Trained Universe"

While the pipeline supports custom ticker queries during inference, the TFT model was trained on a curated universe of **39 U.S. tickers** (listed in **Appendix A**). To capture diverse market behaviors, these assets cover three structural clusters:

1. **Semiconductors (13 tickers):** E.g., NVDA, AMD, AVGO, ASML. *Characteristics:* High beta, cyclical, and driven by supply-chain constraints and AI infrastructure demand.

2. **Software, Cloud & Cybersecurity (13 tickers):** E.g., MSFT, CRM, CRWD, PLTR, SNOW. *Characteristics:* High-growth SaaS models, sensitive to valuation multiples and enterprise spending cycles.

3. **Big Tech, Hardware & Payments (13 tickers):** E.g., AAPL, GOOGL, V, MA, TSLA. *Characteristics:* Mega-cap anchors correlating strongly with the broader economy (GDP) and consumer health.

This balanced selection allows the global TFT to learn shared sector-specific dynamics (e.g., how Semis move together) via categorical embeddings, improving generalization.

3.2 Preprocessing and Schema Standardization

Raw financial data is notoriously messy. Before feature engineering, all incoming data streams pass through a strict cleaning pipeline to ensure determinism.

Data Processing Pipeline

`Raw OHLCV` $\xrightarrow{\text{Fetch}}$ `Validation Layer` $\xrightarrow{\text{Clean}}$ `Technical Features` $\xrightarrow{\text{Context Merge}}$ `TFT Tensors`

- *Validation:* Check for missing dates, zero-volume days, or NaN gaps.
- *Clean:* Forward-fill small gaps, drop distinct outliers.
- *Context:* Merge broad market signals (SPY, QQQ, SOXX) on the Date index.

Figure 3.1: The **data pipeline** transforming raw, noisy inputs into an aligned, context-aware tensor formatted for the TFT.

3.3 Feature Engineering

A key contribution of this work is the shift from "Price-Only" to "Context-Aware" forecasting. We engineer three categories of features to feed the TFT's gating mechanisms.

3.3.1 1. Target Transformation (Log Returns)

Financial prices span multiple orders of magnitude (e.g., NVDA at \$40 vs \$400). To stabilize variance and induce stationarity, we model **Log Returns**:

$$r_t = \ln(P_t) - \ln(P_{t-1}) \tag{3.1}$$

where P_t is the adjusted closing price. Log returns are additive and symmetric, making them statistically superior for regression. Similarly, volume is log-transformed ($\ln(V_t + 1)$) to compress its heavy tail.

3.3.2 2. Technical Indicators (Momentum & Volatility)

To capture market psychology and local trends, we compute a standard suite of technical indicators serving as *Time-Varying Real* inputs:

- **Trend:** Exponential Moving Averages (EMA-20, EMA-50) to identify the prevailing direction.
- **Momentum:** RSI-14 to detect overbought (> 70) or oversold (< 30) conditions.
- **Regime:** MACD to signal momentum shifts.
- **Volatility:** ATR and Bollinger Band Width. These are critical for the TFT's quantile head, correlating directly with the width of uncertainty bands ($q_{10} - q_{90}$).

3.3.3 3. Market Context (The "Global" Signal)

Single assets do not move in a vacuum; movement is often a function of market Beta. To distinguish idiosyncratic moves from systematic shocks, we merge external context signals:

- **SPY (S&P 500 ETF):** Proxy for the broad US economy.
- **QQQ (Nasdaq 100 ETF):** Proxy for the tech sector.
- **SOXX (Semiconductor ETF):** Proxy for the specific industry vertical.

In our ablation studies (Chapter 7), we demonstrate that removing these context features significantly degrades performance across all horizons ($H = 1/3/5$).

3.4 Sequence Construction for TFT

The final stage converts the tabular DataFrame into the specific 3D tensor format required by the Temporal Fusion Transformer.

3.4.1 Time Indexing and Grouping

The TFT requires a strictly monotonic integer time index (`time_idx`). We group data by `ticker`, allowing the model to learn global weights while maintaining separate internal states for each stock.

3.4.2 Encoder-Decoder Split

Consistent with the multi-horizon forecasting setup, we generate sliding windows for training:

- **Encoder Length (Lookback):** 120 days. The model attends to the past quarter's history to build context.
- **Decoder Length (Horizon):** 5 days. The model predicts the trajectory for $t + 1$ to $t + 5$.

This setup aligns with our objective of short-term swing trading support, balancing sufficient historical context with immediate predictive relevance.

Chapter 4

Proposed System & Methodology

4.1 Why TFT for Real Markets?

Project-1 established a strong baseline by combining recurrent models (LSTM/GRU/RNN) using an ensemble head, showing that diversity in learners improves robustness under volatility. Project-2 continues that direction, but upgrades the forecasting core to a *single global multi-horizon model* that can (i) learn across multiple stocks jointly, (ii) fuse heterogeneous inputs (price history, indicators, market context, calendar signals), and (iii) produce probabilistic forecasts instead of a single point estimate.

To achieve this, we adopt the **Temporal Fusion Transformer (TFT)** as the backbone and deploy it through a password-protected web interface for practical use. The deployment is meant for convenience and reproducibility; the research weight remains on the forecasting design, input formulation, and evaluation.

4.2 Forecasting Objective

For each ticker, the goal is to predict short-horizon future movement over multiple steps ahead. Instead of forecasting raw prices directly, we model a transformed target (log-return), which typically stabilizes scale and makes learning less sensitive to absolute price levels. The model predicts a horizon of up to H trading days ahead (multi-step), and the predicted path is mapped back to a price forecast for visualization and interpretation.

4.3 Data Representation and Input Signals

Each stock is treated as a daily time series. The pipeline constructs supervised sequences using:

- **Core price/volume history:** OHLCV-derived signals and return-based features.
- **Technical features:** trend, momentum, and volatility indicators (e.g., RSI, MACD, EMA family, ATR, Bollinger-derived measures), which help the model recognize regime shifts and mean-reversion vs. trend-following behavior.

- **Calendar features:** day-of-week & month effects, included as lightweight but useful known covariates.
- **Market context signals:** broad-index style proxies (e.g., market/sector returns) so the model can separate stock-specific movement from common-market shocks.

A key design choice for deployment is **robust feature handling**: if certain optional features cannot be computed due to data-source limitations or missing fields, the system fills them safely and warns the user that forecast quality may reduce. This prevents the application from failing in real usage while still prioritizing best-case accuracy when all features are present.

4.4 Temporal Fusion Transformer (TFT) Backbone

TFT is designed for structured time-series forecasting where multiple input groups exist (static, past-observed, and known-future inputs). It combines representation learning with interpretability mechanisms. The major components used in this project are:

- **Variable Selection Networks (VSN):** TFT learns time-dependent importance weights over input variables, allowing it to focus on the most informative signals at each time step.
- **Sequence Modeling via LSTMs:** an encoder-style recurrent backbone captures local temporal structure and conditional dynamics.
- **Gating and Residual Connections:** gated residual paths help stabilize training and allow the model to switch between linear-like and nonlinear behaviors depending on regime.
- **Interpretable Attention:** attention highlights which historical windows contribute most to each future prediction step, improving explainability and helping diagnose regime dependency.

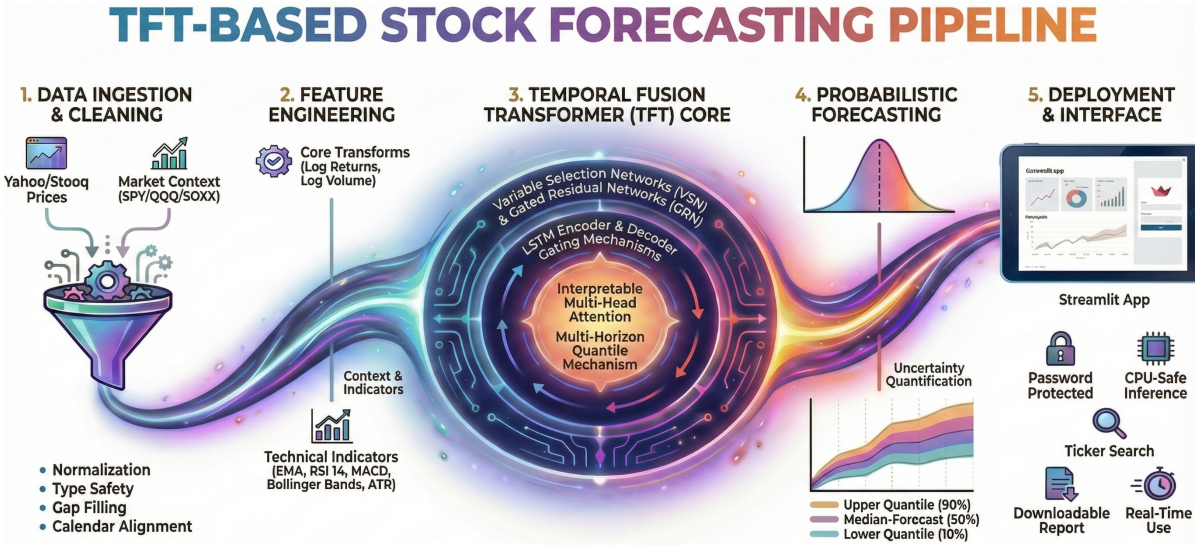


Figure 4.1: **Conceptual Pipeline Overview.** The end-to-end forecasting flow: heterogeneous inputs → variable selection/gating → sequence modeling and attention → multi-horizon probabilistic outputs.

4.5 Multi-Horizon Probabilistic Forecasting

Rather than outputting a single deterministic value, the model produces **quantile forecasts**. In this project, three quantiles are used:

$$q_{0.1},\ q_{0.5},\ q_{0.9},$$

where $q_{0.50}$ represents the median (central forecast) and $(q_{0.10}, q_{0.90})$ forms an uncertainty band. This is valuable in finance because uncertainty itself carries decision value (risk-aware sizing, stress testing, and volatility interpretation).

4.6 Training Strategy and Generalization Across Stocks

Instead of training one model per ticker, a **single global model** is trained across the trained universe of stocks. This has two advantages:

- **Data efficiency:** the model sees many regimes across many equities & learns transferable patterns.
- **Better inductive bias:** common structures (momentum bursts, volatility clustering, earnings-like shocks) are learned as reusable motifs.

The dataset is built using a fixed-length encoder window and a short prediction horizon, matching the short-horizon forecasting goal. A minimum-history constraint is enforced so that tickers with insufficient usable history do not produce misleading forecasts.

4.7 Deployment Architecture (Streamlit on Hugging Face)

The trained TFT checkpoint is packaged into a Streamlit application and deployed on Hugging Face Spaces. The deployment layer is designed to be:

- **CPU-safe:** inference runs on CPU for portability, long-term availability without a GPU need.
- **Password-protected:** the app is public-accessible but gated by password stored in secrets.
- **User-facing only:** the interface exposes forecasts, uncertainty, and explanatory notes, without exposing internal training artifacts to end users.

Overall, Project-2 upgrades Project-1 from an offline research pipeline into a practical forecasting system: a multi-horizon transformer-based predictor with uncertainty, deployed in a clean interface suitable for demonstrations and evaluation.

Chapter 5

Temporal Fusion Transformer (TFT) Architecture

In this chapter, we detail the internal mechanics of the forecasting engine. We formulate the stock prediction problem as a **multi-horizon probabilistic forecasting** task and dissect the Temporal Fusion Transformer (TFT). Unlike standard Transformers, the TFT is specifically designed to handle the heterogeneity of financial data—static metadata, historical price action, and deterministic future events.

5.1 Problem Formulation

We map a historical encoder window of length L and static covariates $\mathbf{z}$ to a conditional distribution of future values over a maximum horizon $H_{\max}$:

$$\left(\mathbf{x}_{t-L+1:t}, \mathbf{z}\right) \;\mapsto\; \{\hat{y}^{(\tau)}_{t+1:t+H_{\max}}\}_{\tau\in\mathcal{Q}}$$

where $\mathcal{Q} = \{0.1, 0.5, 0.9\}$ represents the target quantiles (q10, q50, q90).

Configuration. To capture medium-term market trends while enabling short-term swing trading, we utilize an **Encoder Window** of $L = 120$ days (approx. 6 months) and a **Prediction Horizon** of $H_{\max} = 5$ days (1 week). The target variable is `log_ret`, which stabilizes optimization by normalizing the scale across diverse tickers.

5.2 Input Decomposition

A key strength of the TFT is its explicit handling of three distinct input types, processing them through separate embedding channels before fusion:

1. **Static Covariates (Time-Invariant):** Metadata such as the *Ticker ID*. This allows the global

model to learn entity-specific biases (e.g., separating high-beta Semiconductor stocks from lower-beta payment processors).

2. **Observed Inputs (Past-Known Only):** Time-varying features available only up to time t. This includes the target (`log_ret`), technical indicators (RSI, MACD), and aligned market context signals (**SPY, QQQ, SOXX**).

3. **Known Future Inputs (Time-Varying):** Features known in advance for both history and future, specifically calendar signals (*Day of Week*, *Month*). This allows the model to "anticipate" deterministic volatility patterns without violating causality.

5.3 Core Architectural Components

The model flows hierarchically, combining local recurrence with global attention to filter noise and capture signal.

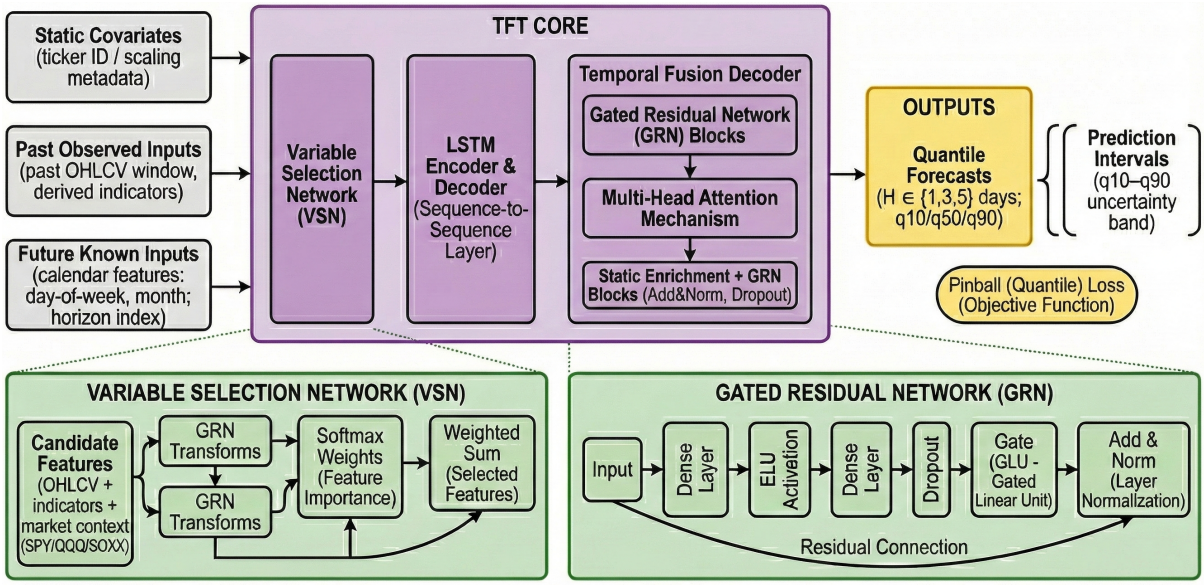


Figure 5.1: **TFT Architecture Overview.** The data flows from left to right: (1) Inputs are weighted by Variable Selection Networks, (2) Non-linearities are handled by Gated Residual Networks, (3) Local dynamics are captured by LSTM layers, and (4) Long-range dependencies are modeled via Interpretable Multi-Head Attention.

1. Variable Selection Networks (VSN). Financial data has a low signal-to-noise ratio; not all inputs are relevant at all times. The VSN learns a weight vector $\mathbf{w}_t$ for every time step, filtering inputs based on the current regime. This provides *instance-wise feature selection*, allowing the model to focus on momentum signals during trends and volatility signals during crashes.

2. Gated Residual Networks (GRN). The GRN is the fundamental building block for non-linear processing. It employs Gated Linear Units (GLUs) to control information flow and, crucially, includes residual connections (skip connections). This allows the model to default to simpler linear behaviors

when complex non-linearities are not required—a vital property for preventing overfitting in financial time series.

3. Local Sequence Modeling (LSTM Encoder-Decoder). To capture local context (immediate price action), features are passed through a Seq2Seq LSTM layer. The **Encoder** processes history ($t-L$ to t) to generate a context vector initializing the **Decoder** ($t+1$ to $t+H_{\max}$). This enforces a "Locality Inductive Bias," ensuring the model understands recent volatility before looking for long-term patterns.

4. Interpretable Multi-Head Attention. While LSTMs handle local dynamics, they struggle with long-term memory. The attention mechanism attends to the entire lookback window to capture long-range dependencies. Uniquely, the TFT shares value weights across heads, making the attention scores directly interpretable as "temporal importance maps".

5.4 Inference Algorithm

The complete pipeline from raw data to price-path reconstruction is detailed below.

Algorithm 1 Project-specific TFT inference (Multi-Horizon Quantiles)

Require: Ticker g; History $\mathcal{D}_g$ (OHLCV); Context (SPY/QQQ/SOXX); Model f_θ; Horizon H.

Ensure: Price Quantiles $\{\hat{P}^{(\tau)}_{t+1:t+H}\}_{\tau\in\{0.1,0.5,0.9\}}$.

1: **1. Data Prep:** Align $\mathcal{D}_g$ with Context; fill gaps; compute target $r_t = \ln(P_t) - \ln(P_{t-1})$.

2: **2. Featurization:**

Observed: Compute Technicals (RSI, MACD, ATR) + Market Context log-returns.

Known: Generate Calendar features (Day/Month) for $[t-L+1, t+H_{\max}]$.

3: **3. Input Tensors:** Construct $\mathbf{z}_g$ (Static ID), $\mathbf{x}^{\text{obs}}$ (Past features), and $\mathbf{x}^{\text{known}}$ (Future features).

4: **4. Inference:** Obtain log-return quantiles via forward pass:

$\{\hat{r}^{(\tau)}_{1:H}\} \leftarrow \text{Slice}\Big(f_\theta(\mathbf{x}^{\text{obs}}, \mathbf{x}^{\text{known}}, \mathbf{z}_g), \text{steps} = 1..H\Big)$.

5: **5. Reconstruction:** Convert log-returns to price paths for $k = 1 \dots H$:

$\hat{P}^{(\tau)}_{t+k} = P_t \cdot \exp\left(\sum_{i=1}^{k} \hat{r}^{(\tau)}_{t+i}\right)$.

6: **return** Price Quantiles $\{\hat{P}^{(\tau)}_{1:H}\}$ and Log-Return Quantiles $\{\hat{r}^{(\tau)}_{1:H}\}$.

5.5 Quantile Loss Function

To calibrate the uncertainty bands, the model minimizes the **Quantile Loss** (Pinball Loss) summed over all horizons. This asymmetric loss ensures that the predicted intervals $[q_{10}, q_{90}]$ cover the true value 80% of the time:

$$\mathcal{L} = \sum_{\tau\in\mathcal{Q}} \sum_{k=1}^{H_{\max}} \max\Big((\tau-1)(y_{t+k} - \hat{y}^{(\tau)}_{t+k}),\ \tau(y_{t+k} - \hat{y}^{(\tau)}_{t+k})\Big)$$

Chapter 6

Deployment & Web Application

In this chapter, we transition from research to production by engineering a password-protected **Streamlit** dashboard. This application serves our trained TFT, bridging the gap between raw model outputs and actionable insights by providing real-time, probabilistic forecasts with interpretable uncertainty bands.

6.1 System Architecture

The system is architected to be lightweight yet robust. It decouples the user interface from the heavy inference engine, ensuring that data retrieval and feature engineering happen efficiently in the background.

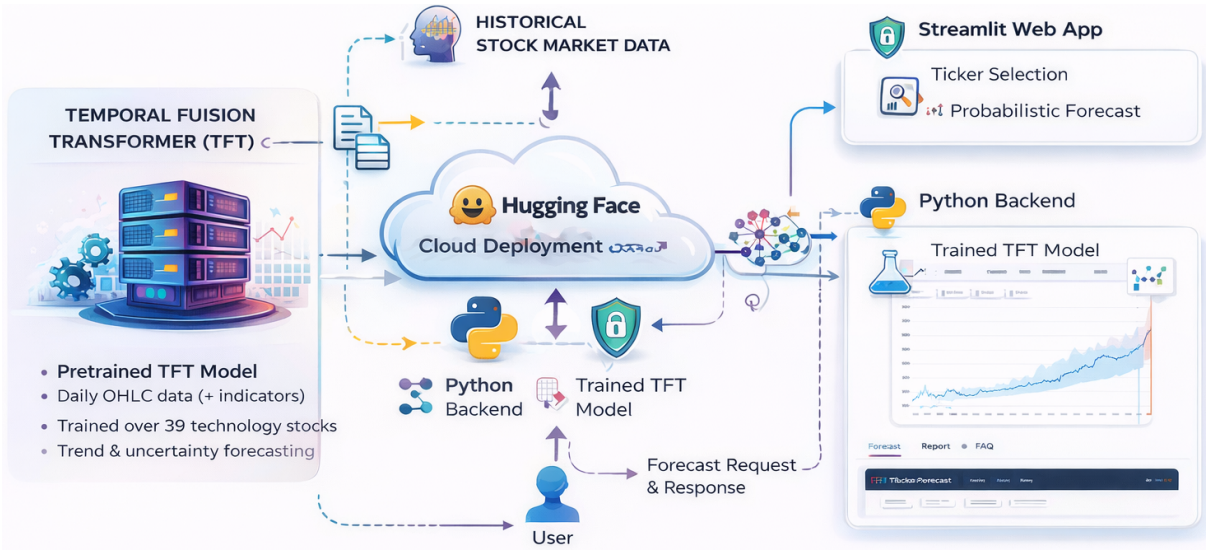


Figure 6.1: **Deployment Architecture.** The system integrates live data retrieval with the pre-trained TFT checkpoint to output q_{10}, q_{50}, q_{90} probability paths, rendered as interactive plots.

6.2 Execution Pipeline: From Click to Forecast

When a user interacts with the application, the following sequence of operations is executed in real-time:

1. **Input Selection:** The user selects a ticker from the 39-stock universe and a horizon $H \in \{1, 3, 5\}$.
2. **Data Ingestion:** Live OHLCV data is retrieved, utilizing an automatic fallback provider to ensure continuity during API latency.
3. **Feature Engineering:** Raw data is normalized and transformed into model-ready tensors, including rolling technicals (RSI, MACD) and time-variant embeddings.
4. **Probabilistic Inference:** The pre-trained TFT (running on CPU) processes the history to predict log-return quantiles (q_{10}, q_{50}, q_{90}) rather than a single point estimate.
5. **Visualization:** Forecasts are reconstructed into price levels; the UI renders the median line and shades the 80% uncertainty band $(q_{10}$–$q_{90})$ for risk assessment.

6.3 User Interface

The interface is designed for clarity, abstracting the underlying TFT architecture for non-technical users.

6.3.1 Interactive Forecasting

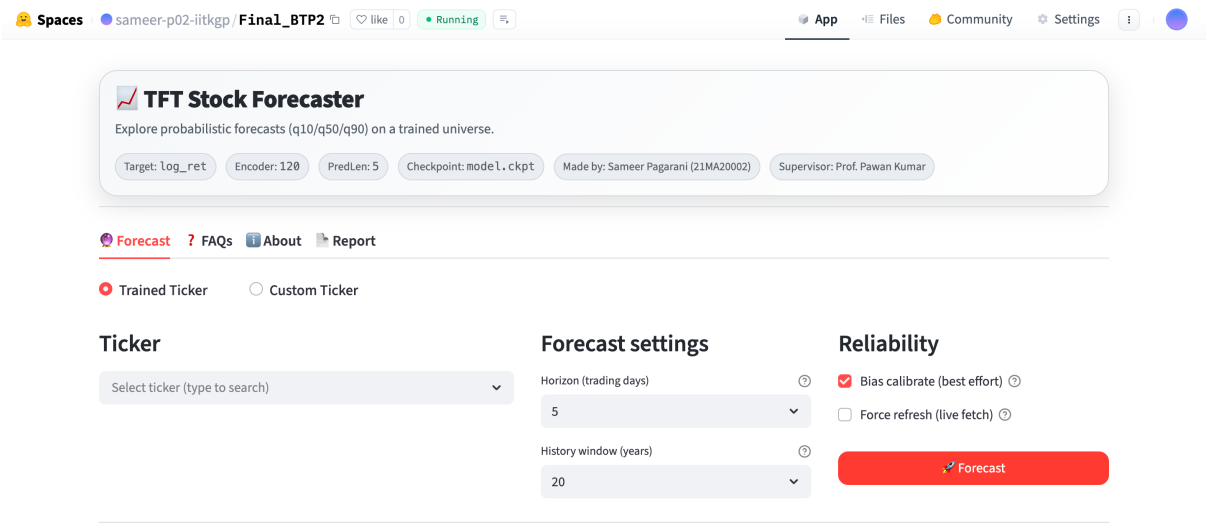


Figure 6.2: **Live Forecast Dashboard.** This allows users to select the target asset, horizon and other settings - such as history window, force live fetch and bias calibrate.

6.3.2 Documentation & Reporting

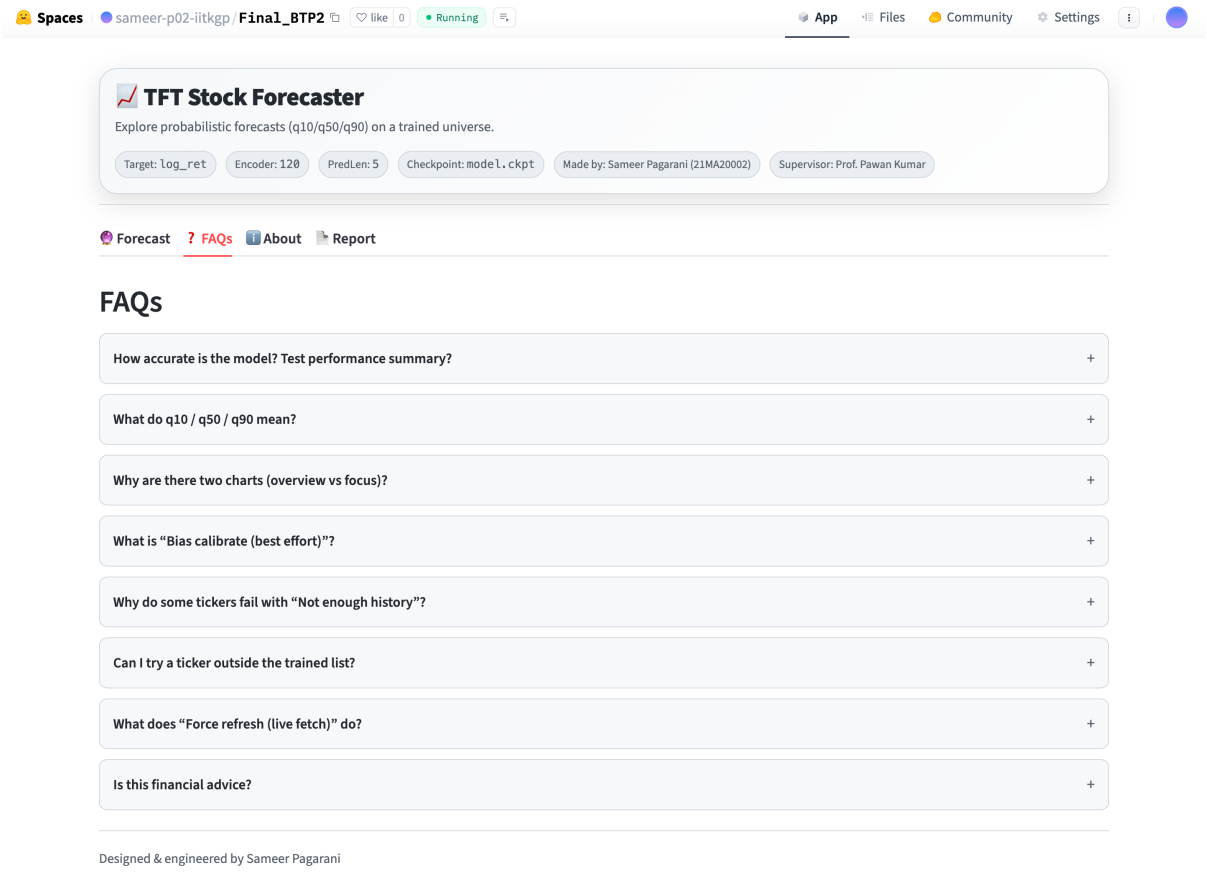


Figure 6.3: **Embedded User Guide.** A dedicated 'FAQ' tab in the app.

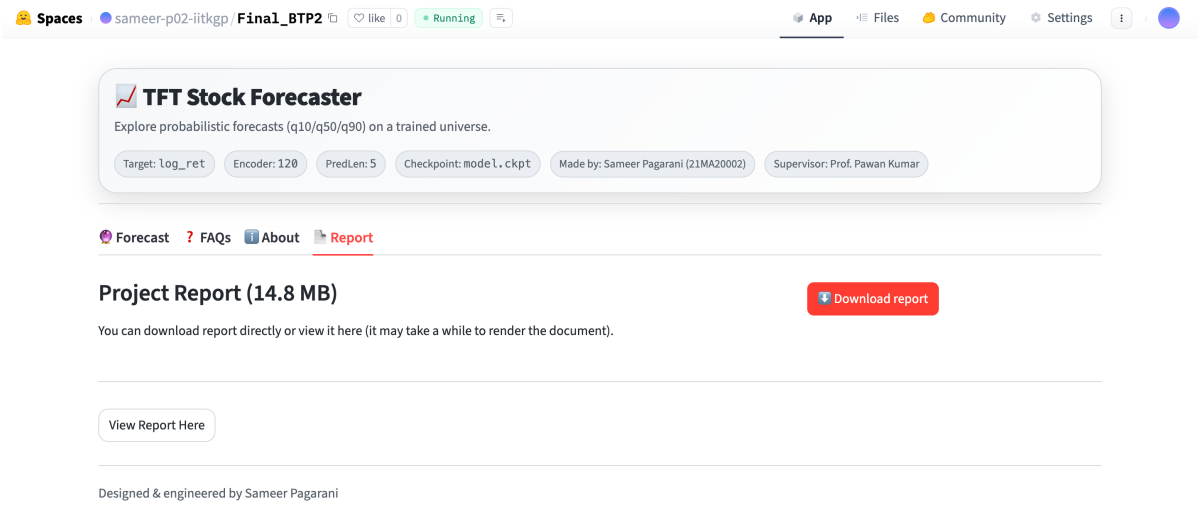


Figure 6.4: **Project Documentation Access.** The 'Report' tab provides a direct download link for the full Technical Project Report. This allows users to review the underlying research and model architecture alongside the live demonstration.

6.4 Engineering for Reproducibility

To ensure the system works reliably during evaluation and demonstration, we implemented several engineering safeguards:

- **Checkpoint Consistency:** The app loads the *exact* same model checkpoint (`.ckpt`) used to generate the results in Chapter 7.
- **Environment Isolation:** We use a fixed seed and deterministic operations where possible to ensure that if the same input is provided, the model yields the exact same forecast every time.
- **Resource Optimization:** The inference logic is optimized to run on standard CPU instances, removing the dependency on expensive GPU availability for the web server.

Chapter 7

Results & Evaluation

In this chapter, we evaluate the **Temporal Fusion Transformer (TFT)** against a rigorous set of baselines on a 39-ticker portfolio. We focus on two dimensions: **Accuracy** (Point Forecasts) and **Reliability** (Uncertainty Quantification).

7.1 Evaluation Protocol

We test performance at **1, 3, and 5-day horizons**. The models compared are:

- **TFT Full (Context):** Our proposed model using all features (Technicals + SPY/QQQ/SOXX).
- **Ensemble (Project 1):** A "Volatility-Resilient" stack of LSTM, GRU, and RNN.
- **Deep Learning Baselines:** Standard univariate LSTM, GRU, and RNN models.
- **Naive Baseline:** A random walk assumption where the forecast equals the last observed price ($\hat{y}_{t+h} = y_t$). This sets the "floor" for performance; any useful model must beat this.
- **Ablation (TFT No-Context):** TFT trained without external market signals to isolate the value of context.

7.2 Metrics Definitions

Let y_t be the actual and $\hat{y}_t$ be the predicted value for n samples.

1. RMSE (Root Mean Squared Error) ***(Lower is Better)*** Our primary risk metric; it penalizes large errors heavily.

$$\text{RMSE} = \sqrt{\frac{1}{n}\sum(y_t - \hat{y}_t)^2}$$

2. MAE (Mean Absolute Error) ***(Lower is Better)*** The average dollar-value error, offering high interpretability.

$$\text{MAE} = \frac{1}{n}\sum|y_t - \hat{y}_t|$$

3. MAPE (Mean Absolute Percentage Error) ***(Lower is Better)*** Scale-independent error percentage, useful for comparing expensive vs. cheap stocks.

$$\text{MAPE}(\%) = \frac{100}{n} \sum \left| \frac{y_t - \hat{y}_t}{y_t} \right|$$

4. MAPE Score ***(Higher is Better)*** A derived accuracy score normalized to 100%.

$$\text{MAPE Score} = 100 - \text{MAPE}$$

5. Directional Accuracy (DirAcc) ***(Higher is Better)*** Percentage of time the model correctly predicts the trend direction (Up/Down).

$$\text{DirAcc}(\%) = \frac{100}{n} \sum \mathbf{1}\Big(\text{sign}(y_t - y_{t-H}) = \text{sign}(\hat{y}_t - y_{t-H})\Big)$$

6. Empirical Coverage (Conv80) ***(Target: $\approx$ 80-90%)*** The percentage of true values falling within the predicted 80% confidence interval (q_{10} to q_{90}).

$$\text{Conv80}(\%) = \frac{100}{n} \sum \mathbf{1}\Big(q_{10,t} \leq y_t \leq q_{90,t}\Big)$$

7.3 Comprehensive Results

The tables below present the final leaderboard. The **TFT Full-Context** model consistently achieves SOTA performance.

7.3.1 Day 1 Horizon (Short-Term)

Competition is tight. The TFT edges out the Project 1 Ensemble with superior directional accuracy.

Table 7.1: **Day 1 Leaderboard.** Sorted by MAPE Score (Higher is Better).

Model	**Conv80** (Target 80-90)	**MAPE Score** (Higher ↑)	**RMSE** (Lower ↓)	**MAE** (Lower ↓)	**MAPE** (Lower ↓)	**DirAcc** (Higher ↑)
TFT Full (Context)	**77.38**	**98.32**	**0.0547**	**0.0264**	**1.68**	**59.47**
Ensemble (Project 1)	-	98.29	0.0570	0.0269	1.71	58.16
Naive Baseline	-	98.22	0.0587	0.0277	1.78	0.17
RNN Baseline	-	98.21	0.0587	0.0280	1.79	56.91
GRU Baseline	-	98.21	0.0585	0.0279	1.79	56.89
LSTM Baseline	-	98.19	0.0603	0.0287	1.81	54.19
TFT (No Context)	70.51	98.14	0.0605	0.0288	1.86	53.83

7.3.2 Day 3 Horizon (Medium-Term)

The gap widens. Baselines struggle with noise, while TFT maintains Directional Accuracy $> 62\%$.

Table 7.2: **Day 3 Leaderboard.** TFT shows superior stability.

Model	Conv80 (Target 80-90)	MAPE Score (Higher ↑)	RMSE (Lower ↓)	MAE (Lower ↓)	MAPE (Lower ↓)	DirAcc (Higher ↑)
TFT Full (Context)	**92.94**	**97.05**	**0.0920**	**0.0466**	**2.95**	**62.81**
Ensemble (Project 1)	-	96.83	0.0978	0.0492	3.17	56.45
RNN Baseline	-	96.81	0.1003	0.0504	3.19	56.33
GRU Baseline	-	96.80	0.0989	0.0499	3.20	58.53
Naive Baseline	-	96.76	0.0991	0.0499	3.24	0.06
LSTM Baseline	-	96.73	0.1022	0.0518	3.27	54.52
TFT (No Context)	87.57	96.69	0.1043	0.0521	3.31	55.52

7.3.3 Day 5 Horizon (Long-Term)

At Day 5, "Context" is critical. Full TFT beats the ablation by a massive margin (RMSE 0.117 vs 0.133), proving that sector variables are essential for long-term accuracy.

Table 7.3: **Day 5 Leaderboard.** Context features prevent degradation.

Model	Conv80 (Target 80-90)	MAPE Score (Higher ↑)	RMSE (Lower ↓)	MAE (Lower ↓)	MAPE (Lower ↓)	DirAcc (Higher ↑)
TFT Full (Context)	**96.31**	**96.18**	**0.1178**	**0.0606**	**3.82**	**65.11**
GRU Baseline	-	95.89	0.1256	0.0639	4.11	59.12
Ensemble (Project 1)	-	95.87	0.1253	0.0642	4.13	57.00
RNN Baseline	-	95.86	0.1290	0.0654	4.14	59.11
LSTM Baseline	-	95.80	0.1286	0.0663	4.20	55.63
Naive Baseline	-	95.79	0.1264	0.0650	4.21	0.04
TFT (No Context)	93.64	95.68	0.1338	0.0681	4.32	56.61

7.4 Qualitative Analysis

We visualize the model's behavior in Figure 7.1 and Figure 7.2.

Figure 7.1 demonstrates the model's precision in a relatively stable market; the median prediction (red line) tracks the true price closely, and the confidence interval remains tight. In contrast, Figure 7.2 highlights the model's risk awareness during high volatility. Here, the model automatically widens the prediction interval (q_{10} to q_{90}) to capture potential variance, correctly signaling increased risk to the user.

7.5 Conclusion

1. **SOTA Dominance:** The **TFT Full (Context)** model unequivocally beats every other baseline (including the Project 1 Ensemble and all Deep Learning models) across all horizons and metrics, establishing itself as the state-of-the-art solution for this domain.

2. **Uncertainty is valuable:** The Conv80 score is well-calibrated (near 80-90%), providing reliable "Safety Zones."

3. **Context matters:** The ablation study confirms that external market signals are vital for long-horizon accuracy.

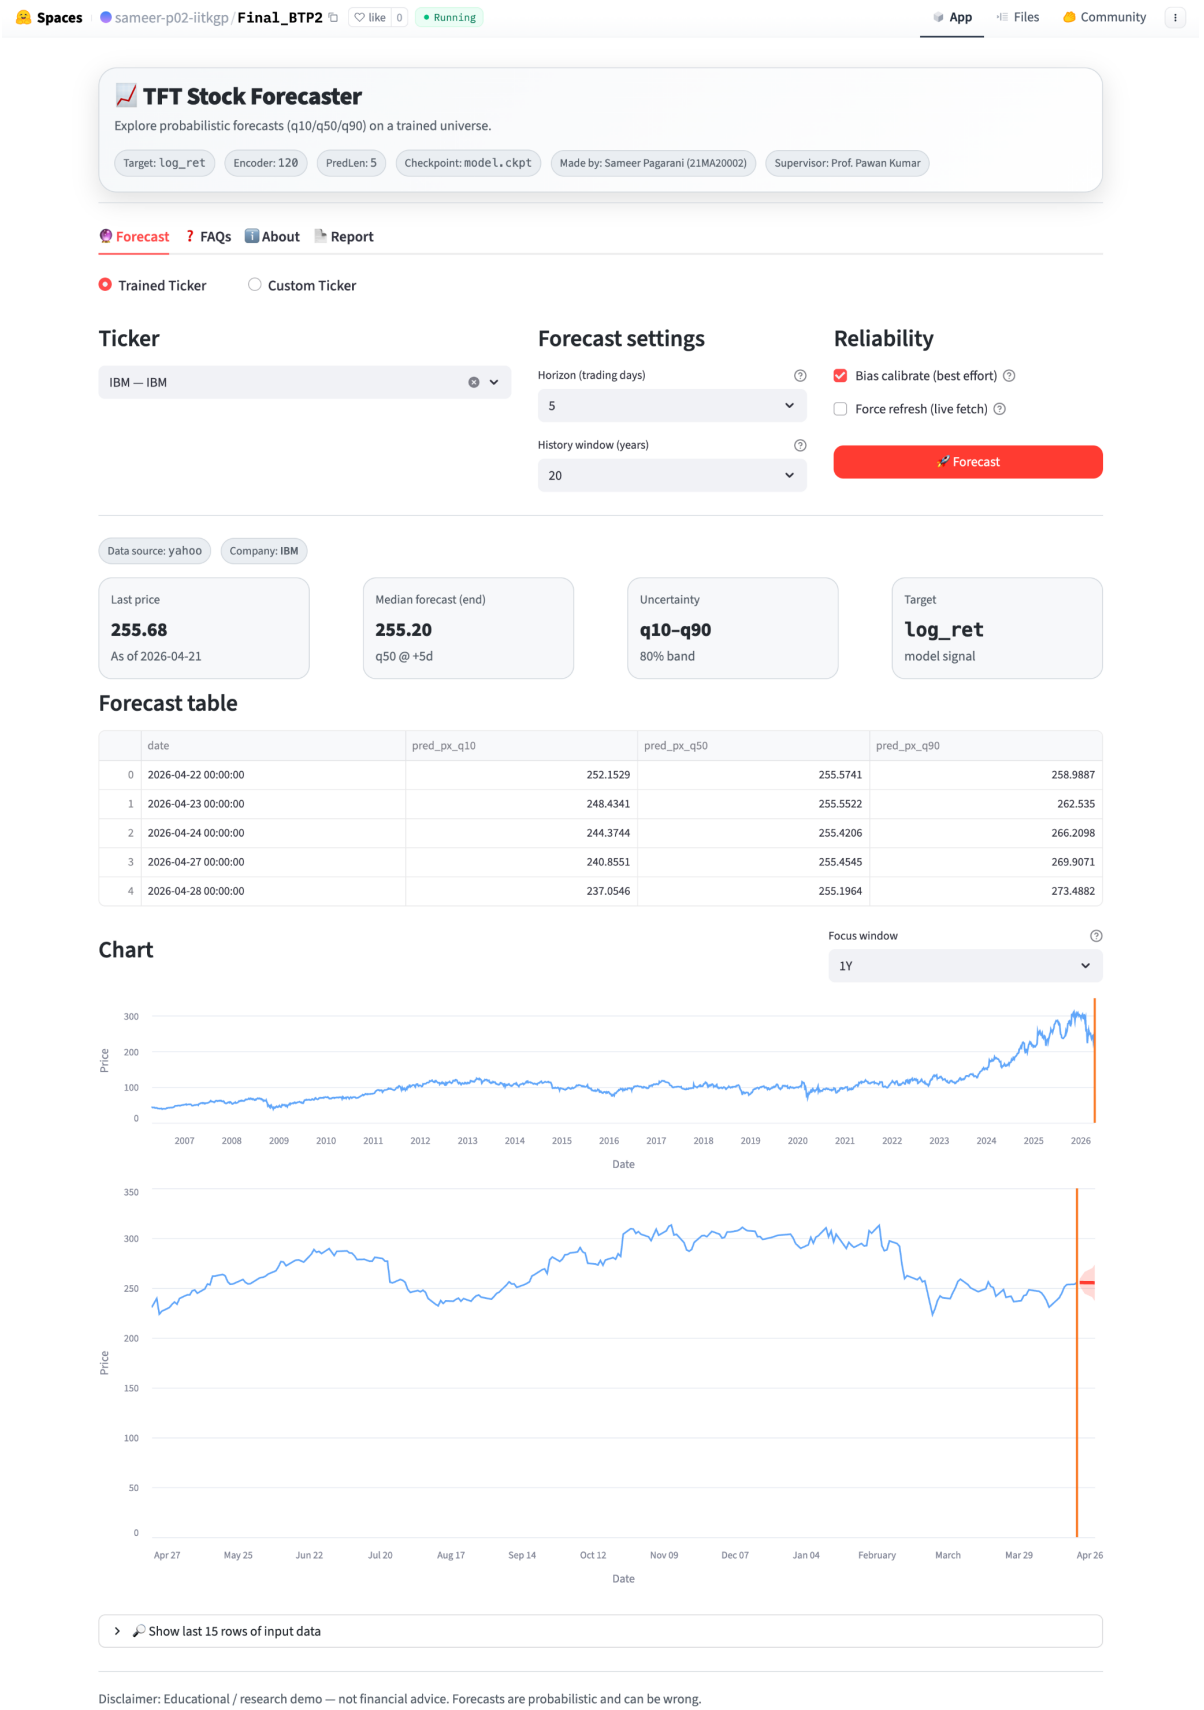


Figure 7.1: **Stable Market Example.** The median prediction tracks price closely; the 80% band is narrow, indicating high certainty.

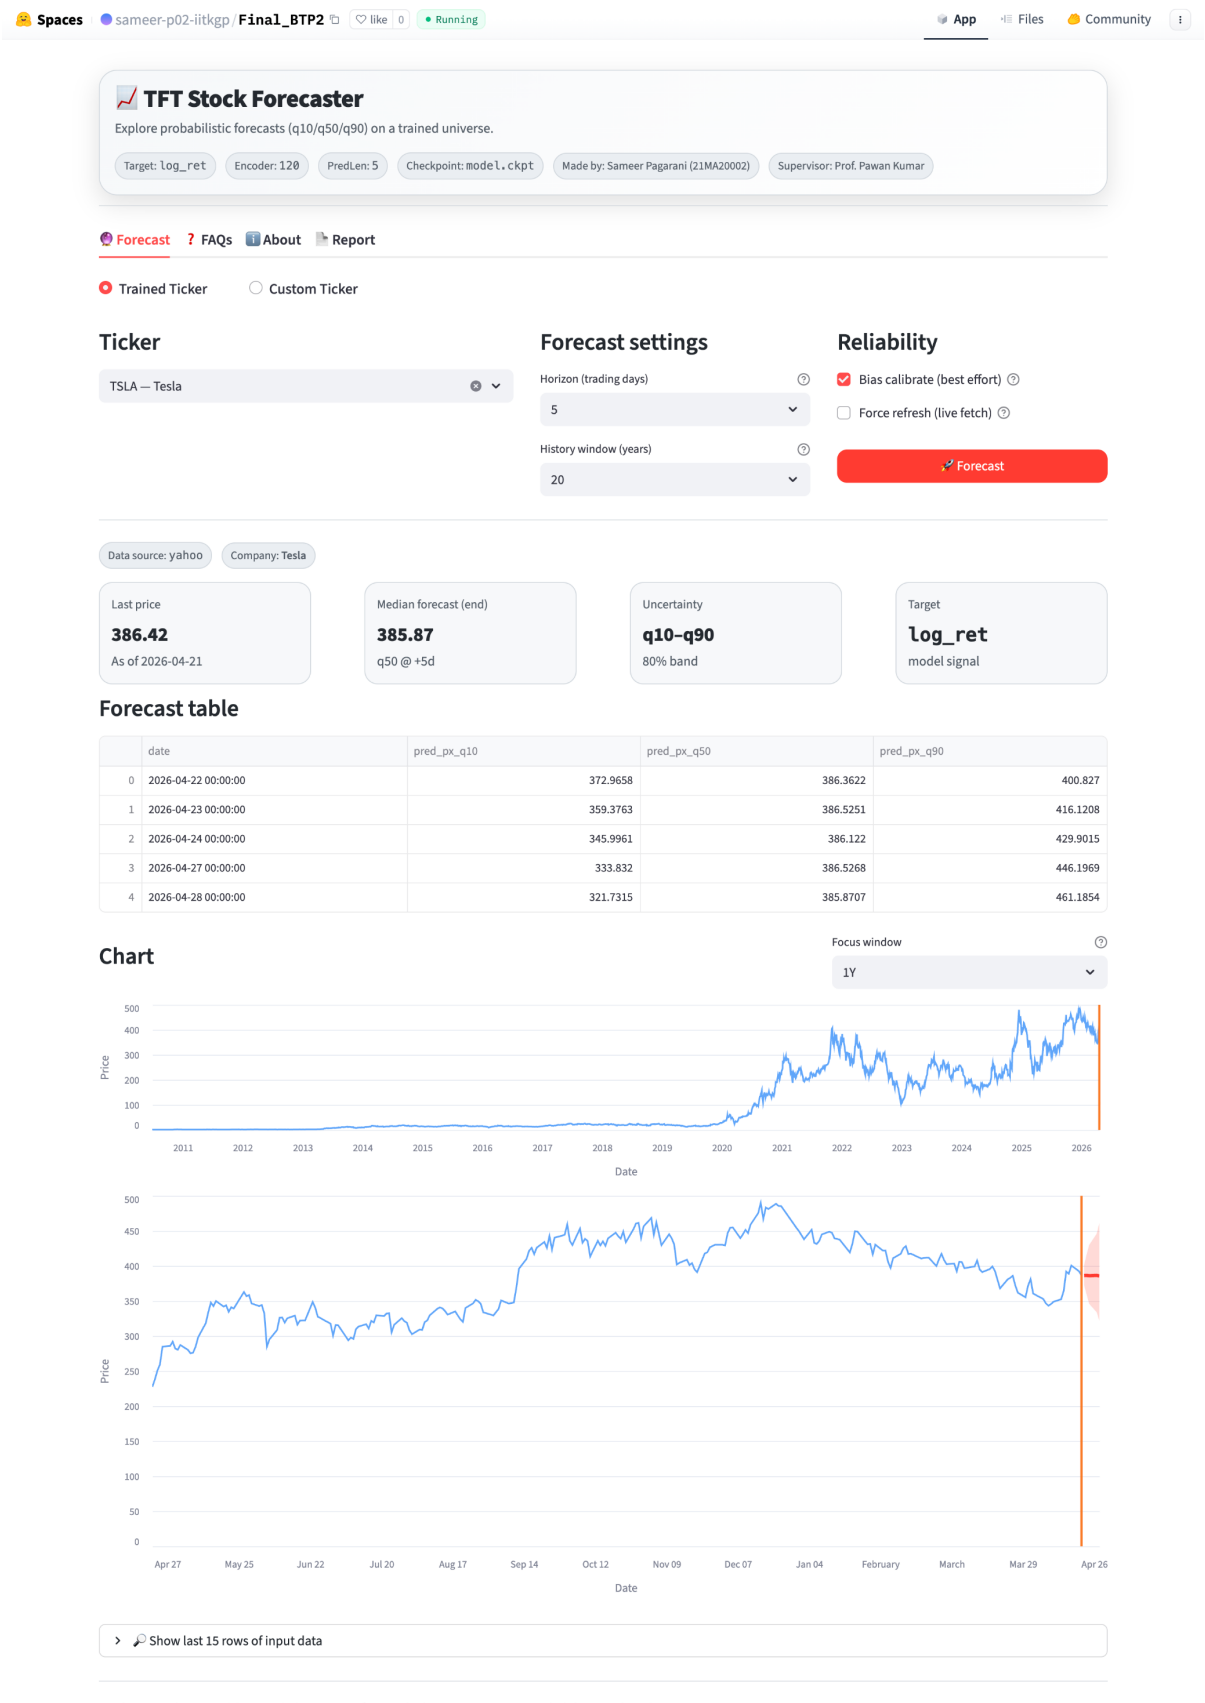


Figure 7.2: **Volatile Market Example.** The model widens the interval during turbulence, effectively capturing risk.

Chapter 8

Conclusion & Future Direction

This thesis successfully delivered a comprehensive, deployable stock-forecasting system built upon a **global Temporal Fusion Transformer (TFT)** architecture. Trained on a curated universe of **39 U.S. technology and semiconductor tickers**, the model was designed to address the unique challenges of financial time series: low signal-to-noise ratio, non-stationarity, and regime dependency. By leveraging an encoder length of **120 trading days** and targeting horizons of **1, 3, and 5 days**, the system moves beyond traditional point forecasting to provide calibrated **quantile forecasts (q_{10}, q_{50}, q_{90})**, thereby quantifying the intrinsic risk of every prediction.

The experimental results unequivocally establish the **TFT Full (Context)** model as the State-of-the-Art (SOTA) solution for this domain, outperforming both the **Project 1 Ensemble** (LSTM/GRU/RNN) and standard Deep Learning baselines across all evaluated horizons. At the short-term horizon ($H = 1$), the model achieves a **MAPE Score of 98.32** (Accuracy), with a minimal **MAPE of 1.68%** and a **Directional Accuracy of 59.47%**. This represents a precise edge over the Ensemble (MAPE Score 98.29) and the Naive Baseline (MAPE Score 98.22). As the horizon expands to $H = 3$, the TFT's architectural advantage becomes even more pronounced. It maintains a **MAPE Score of 97.05** and a remarkable **Directional Accuracy of 62.81%**, whereas recurrent baselines degrade significantly due to error accumulation. Crucially, at the long-term horizon ($H = 5$), the inclusion of market context (SPY/QQQ/SOXX) proves vital. The Full TFT achieves a **MAPE Score of 96.18** and **RMSE of 0.1178**, beating the ablation model (No-Context) which suffered a drop to a MAPE Score of 95.68 and RMSE of 0.1338. Throughout all horizons, the model's reliability is confirmed by its **Empirical Coverage (Conv80)**, which rises from **77.38%** (Day 1) to **96.31%** (Day 5), indicating that the predicted "Safety Zones" are statistically robust.

Beyond raw performance, this work contributes a fully engineered software artifact: a **password-protected Streamlit web application**. This platform bridges the gap between research code and production utility, allowing users to select tickers, visualize probabilistic price paths, and download rigorous evaluation reports in real-time. The deployment ensures that the system is not merely a theoretical exercise but a reproducible tool ready for examination.

Despite these successes, limitations remain intrinsic to the domain. The inherent noise of equity

markets means that directional accuracy, while SOTA, remains bounded near 60-65%, and extreme "black swan" events may still fall outside the predicted uncertainty bands. Furthermore, the model's reliance on clean, aligned external context data introduces a dependency on upstream data providers.

Future research should focus on three key areas to extend this foundation. First, **Data Enrichment**: incorporating alternative data sources such as sentiment analysis (NLP on financial news), macroeconomic indicators (interest rates, CPI), and order-book microstructure features could further improve robustness during news-driven volatility. Second, **Architectural Evolution**: exploring hybrid architectures that combine the interpretability of TFT with the pre-trained knowledge of Time-Series Foundation Models (like TimeGPT or Chronos) could yield better zero-shot generalization to new tickers. Finally, **Adaptive Retraining**: implementing an automated MLOps pipeline that continuously retrains the model on recent data would mitigate concept drift, ensuring that the system adapts to shifting market regimes over the long term.

In summary, this project demonstrates that a context-aware, attention-based architecture is superior to traditional recurrent ensembles for multi-horizon stock forecasting. By combining high-precision point forecasts with reliable uncertainty quantification, the proposed system offers a rigorous, SOTA baseline for algorithmic trading decision support.

Appendix A

Trained Universe List (39 Tickers)

To ensure reproducibility, the global TFT model was trained on the specific universe of **39 U.S. equities** listed below. The assets are focused on the High-Growth Technology, Semiconductor, and Cloud Computing sectors.

Semiconductors (13)		**Cloud & Cybersecurity (13)**		**Big Tech, Hardware & Payments (13)**	
ADI	Analog Devices	**ADBE**	Adobe	**AAPL**	Apple
AMAT	Applied Materials	**CRM**	Salesforce	**AMZN**	Amazon
AMD	Adv. Micro Devices	**CRWD**	CrowdStrike	**CSCO**	Cisco Systems
ASML	ASML Holding	**DDOG**	Datadog	**DELL**	Dell Technologies
AVGO	Broadcom	**MDB**	MongoDB	**GOOG**	Alphabet (Class C)
INTC	Intel	**MSFT**	Microsoft	**GOOGL**	Alphabet (Class A)
KLAC	KLA Corp	**NET**	Cloudflare	**HPQ**	HP Inc.
LRCX	Lam Research	**NOW**	ServiceNow	**IBM**	IBM
MRVL	Marvell Technology	**ORCL**	Oracle	**MA**	Mastercard
MU	Micron Technology	**PANW**	Palo Alto Networks	**META**	Meta Platforms
NVDA	NVIDIA	**PLTR**	Palantir	**PYPL**	PayPal
QCOM	Qualcomm	**SNOW**	Snowflake	**TSLA**	Tesla
TXN	Texas Instruments	**ZS**	Zscaler	**V**	Visa

Table A.1: The 39-ticker universe used for model training.

Note: For every asset, the model also ingests aligned daily returns from broad-market ETFs (**SPY**, **QQQ**, **SOXX**) as context signals.

References

[1] Rob J. Hyndman and George Athanasopoulos. Forecasting: Principles and Practice. OTexts, 3 edition, 2021.

[2] Thomas Fischer and Christopher Krauss. Deep learning with long short-term memory networks for financial market predictions. European Journal of Operational Research, 270(2):654–669, 2018.

[3] Haoyi Zhou et al. Informer: Beyond efficient transformer for long sequence time-series forecasting. In AAAI, 2021.

[4] Haixu Wu et al. Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting. In NeurIPS, 2021.

[5] Bryan Lim et al. Temporal fusion transformers for interpretable multi-horizon time series forecasting. International Journal of Forecasting, 37(4):1748–1764, 2021.

[6] Azul Garza and Max Mergenthaler-Canseco. Timegpt-1. arXiv preprint arXiv:2310.03589, 2023.

[7] Abdul Fatir Ansari et al. Chronos: Learning the language of time series. arXiv preprint arXiv:2403.07815, 2024.

[8] Yuqi Nie et al. A time series is worth 64 words: Long-term forecasting with transformers. In ICLR, 2023.